

Movies Reviews Classification Using BERT

ID	NAME
6338	Ahmed Mohamed Ahmed Ali
6672	Mostafa Mohamed tawfiq
6160	Mazen medhat Farid
6670	Mohamed zayton



Running model (use Bert directly)

With preprocessing

LR=1e-6

```
▶ model = BertClassifier()  
# 512, 256, 128, 64
```

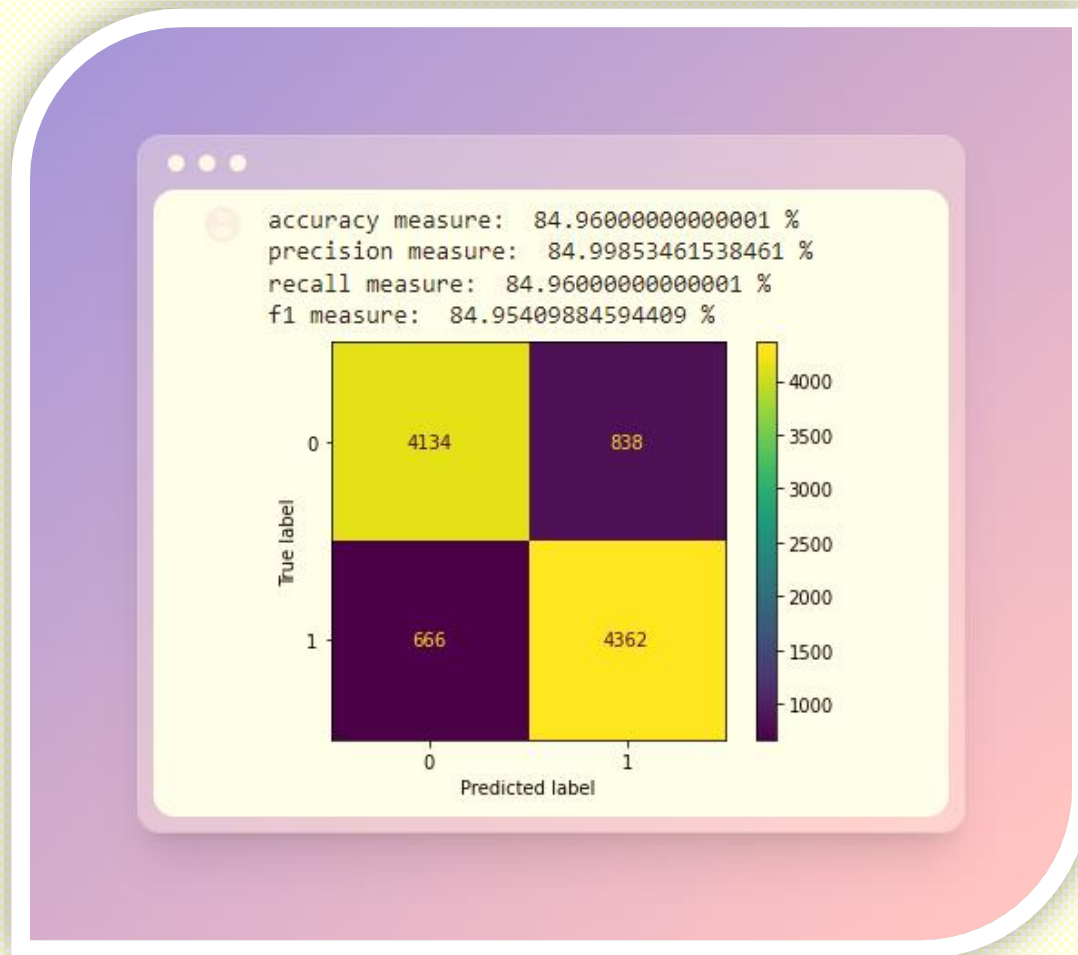
Some weights of the model checkpoint at bert-base-cased were not used when initializing BertModel: ['cls.pr
- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task
- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to

```
[ ] run_model(model, train_dataloader, val_dataloader, LR)
```

```
100%|██████████| 563/563 [17:19<00:00, 1.85s/it]  
100%|██████████| 63/63 [01:47<00:00, 1.70s/it]  
Epochs: 1 | Train Loss: 0.010 | Train Accuracy: 0.583 | Val Loss: 0.008 | Val Accuracy: 0.779  
100%|██████████| 563/563 [16:49<00:00, 1.79s/it]  
100%|██████████| 63/63 [01:47<00:00, 1.71s/it]  
Epochs: 2 | Train Loss: 0.007 | Train Accuracy: 0.784 | Val Loss: 0.007 | Val Accuracy: 0.811  
100%|██████████| 563/563 [16:58<00:00, 1.81s/it]  
100%|██████████| 63/63 [01:48<00:00, 1.72s/it]  
Epochs: 3 | Train Loss: 0.007 | Train Accuracy: 0.805 | Val Loss: 0.006 | Val Accuracy: 0.822  
100%|██████████| 563/563 [17:03<00:00, 1.82s/it]  
100%|██████████| 63/63 [01:45<00:00, 1.68s/it]  
Epochs: 4 | Train Loss: 0.006 | Train Accuracy: 0.820 | Val Loss: 0.006 | Val Accuracy: 0.831  
100%|██████████| 563/563 [17:16<00:00, 1.84s/it]  
100%|██████████| 63/63 [02:32<00:00, 2.42s/it]  
Epochs: 5 | Train Loss: 0.006 | Train Accuracy: 0.829 | Val Loss: 0.006 | Val Accuracy: 0.843  
(29843, 217.27602770924568, 3371, 22.142377495765686)
```

Show confusion matrix

(Best Accuracy)



Running model (using 4 hidden layers)

With preprocessing

LR=1e-6

```
model2 = BertClassifier2(512, 256, 128, 64)
```


Some weights of the model checkpoint at bert-base-cased were not used when initializing BertModel: ['
- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another language (e.g., Spanish) and the vocabulary in the checkpoint is different from the one in this checkpoint.
- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be initialized from this checkpoint.

```
[ ] run_model(model2, train_dataloader, val_dataloader, LR)
```

```
0%|          | 0/563 [00:00<?, ?it/s]/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2
100%|██████████| 563/563 [06:06<00:00, 1.54it/s]
100%|██████████| 63/63 [00:25<00:00, 2.51it/s]
Epochs: 1 | Train Loss: 0.011 | Train Accuracy: 0.531 | Val Loss: 0.011 | Val Accuracy: 0.664
100%|██████████| 563/563 [06:06<00:00, 1.54it/s]
100%|██████████| 63/63 [00:23<00:00, 2.68it/s]
Epochs: 2 | Train Loss: 0.010 | Train Accuracy: 0.740 | Val Loss: 0.010 | Val Accuracy: 0.798
(26627, 368.1936562657356, 3192, 38.11210232973099)
```

Show confusion matrix

```
accuracy measure: 49.72 %  
precision measure: 24.720784 %  
recall measure: 49.72 %  
f1 measure: 33.022687683676196 %
```





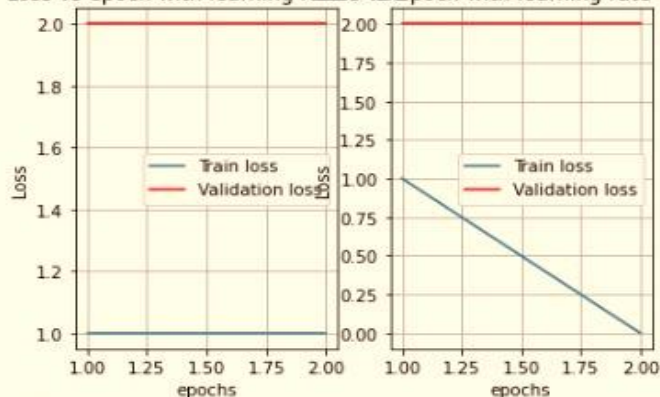
Running model

Without preprocessing

○ LR=1e-2

```
100%|██████████| 563/563 [03:58<00:00, 2.36it/s]
100%|██████████| 63/63 [00:22<00:00, 2.80it/s]
Epochs: 1 | Train Loss: 0.000 | Train Accuracy: 0.500 | Val Loss: 0.001 | Val Accuracy: 0.493
100%|██████████| 563/563 [04:00<00:00, 2.34it/s]
100%|██████████| 63/63 [00:21<00:00, 2.96it/s]
Epochs: 2 | Train Loss: 0.000 | Train Accuracy: 0.500 | Val Loss: 0.001 | Val Accuracy: 0.493
547851
100%|██████████| 563/563 [04:00<00:00, 2.35it/s]
100%|██████████| 63/63 [00:23<00:00, 2.70it/s]
Epochs: 1 | Train Loss: 0.000 | Train Accuracy: 0.500 | Val Loss: 0.001 | Val Accuracy: 0.493
100%|██████████| 563/563 [04:01<00:00, 2.33it/s]
100%|██████████| 63/63 [00:21<00:00, 2.96it/s]
Epochs: 2 | Train Loss: 0.000 | Train Accuracy: 0.500 | Val Loss: 0.001 | Val Accuracy: 0.493
547851
100%|██████████| 563/563 [03:59<00:00, 2.35it/s]
100%|██████████| 63/63 [00:21<00:00, 2.98it/s]
Epochs: 1 | Train Loss: 0.000 | Train Accuracy: 0.500 | Val Loss: 0.001 | Val Accuracy: 0.493
100%|██████████| 563/563 [04:03<00:00, 2.31it/s]
100%|██████████| 63/63 [00:21<00:00, 2.95it/s]
Epochs: 2 | Train Loss: 0.000 | Train Accuracy: 0.500 | Val Loss: 0.001 | Val Accuracy: 0.493
547851
100%|██████████| 563/563 [04:01<00:00, 2.33it/s]
100%|██████████| 63/63 [00:21<00:00, 2.93it/s]
Epochs: 1 | Train Loss: 0.000 | Train Accuracy: 0.500 | Val Loss: 0.001 | Val Accuracy: 0.493
100%|██████████| 563/563 [04:01<00:00, 2.33it/s]
100%|██████████| 63/63 [00:23<00:00, 2.72it/s]
Epochs: 2 | Train Loss: 0.000 | Train Accuracy: 0.500 | Val Loss: 0.001 | Val Accuracy: 0.493
```



Loss vs epoch with learning rate = 0.01



Best loss: 1
Best accuracy: 2
Best learning rate: 1

Show confusion matrix

accuracy measure: 49.72 %
precision measure: 24.720784 %
recall measure: 49.72 %
f1 measure: 33.022687683676196 %






LR=1e-6

```
0%|          | 0/563 [00:00<?, ?it/s]/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py
100%|██████████| 563/563 [05:32<00:00, 1.69it/s]
100%|██████████| 63/63 [00:23<00:00, 2.66it/s]
Epochs: 1 | Train Loss: 0.013 | Train Accuracy: 0.500 | Val Loss: 0.013 | Val Accuracy: 0.507
100%|██████████| 563/563 [05:33<00:00, 1.69it/s]
100%|██████████| 63/63 [00:23<00:00, 2.71it/s]
Epochs: 2 | Train Loss: 0.013 | Train Accuracy: 0.500 | Val Loss: 0.013 | Val Accuracy: 0.507
100%|██████████| 563/563 [05:33<00:00, 1.69it/s]
100%|██████████| 63/63 [00:23<00:00, 2.68it/s]
Epochs: 1 | Train Loss: 0.013 | Train Accuracy: 0.500 | Val Loss: 0.013 | Val Accuracy: 0.507
100%|██████████| 563/563 [05:33<00:00, 1.69it/s]
100%|██████████| 63/63 [00:23<00:00, 2.72it/s]
Epochs: 2 | Train Loss: 0.013 | Train Accuracy: 0.500 | Val Loss: 0.013 | Val Accuracy: 0.507
accuracy measure: 49.72 %
precision measure: 24.720784 %
recall measure: 49.72 %
f1 measure: 33.022687683676196 %
```

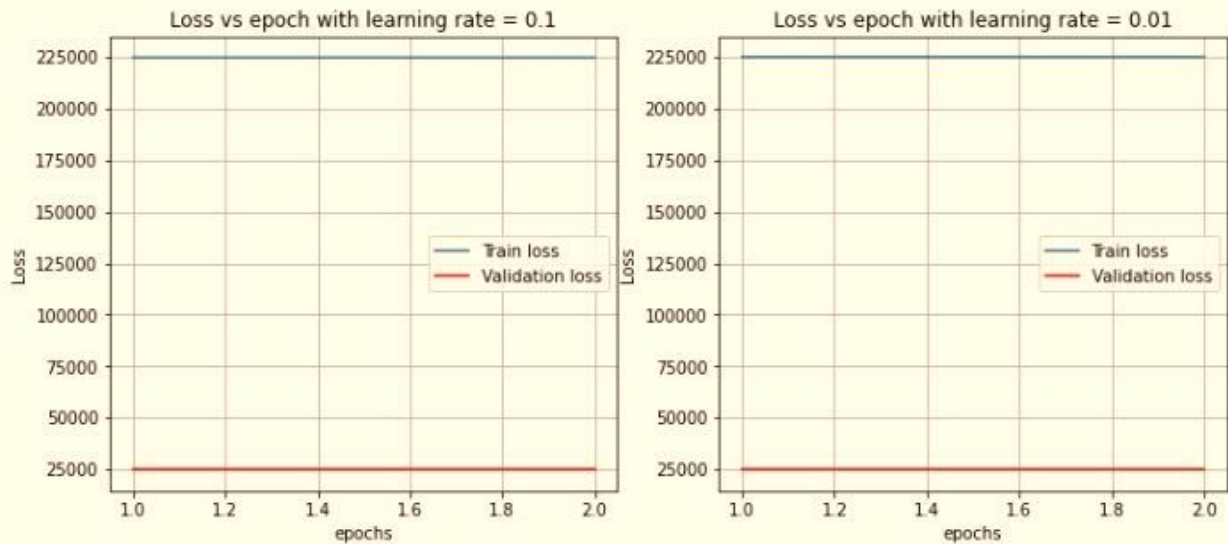

Show confusion matrix

```
accuracy measure: 49.72 %  
precision measure: 24.720784 %  
recall measure: 49.72 %  
f1 measure: 33.022687683676196 %  
/usr/local/lib/python3.7/dist-packages/sklea  
_warn_prf(average, modifier, msg_start, le
```



```
Best loss: 1  
Best accuracy: 2  
Best learning rate: 1
```

comparing between the best validation accuracies for the different values of learning rate:



Best loss: 24662.5
Best accuracy: 1973
Best learning rate: 0.1
Test-acc = 5028

❖ Bonus part hidden layers:

LR = 1e-6


model = BertClassifier3(256, 128, 64, 32, 16)

```
[ ] run_model(model, train_dataloader, val_dataloader, LR)
```

```
0%|          | 0/563 [00:00<?, ?it/s]/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:24:
100%|██████████| 563/563 [06:07<00:00, 1.53it/s]
100%|██████████| 63/63 [00:22<00:00, 2.74it/s]
Epochs: 1 | Train Loss: 0.011 | Train Accuracy: 0.500 | Val Loss: 0.011 | Val Accuracy: 0.493
100%|██████████| 563/563 [06:06<00:00, 1.54it/s]
100%|██████████| 63/63 [00:23<00:00, 2.68it/s]
Epochs: 2 | Train Loss: 0.011 | Train Accuracy: 0.500 | Val Loss: 0.011 | Val Accuracy: 0.493
(17991, 388.54497665166855, 1973, 42.99576020240784)
```

Show confusion matrix

```
accuracy measure: 50.28 %  
precision measure: 25.280784 %  
recall measure: 50.28 %  
f1 measure: 33.64490817141336 %  
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/  
_warn_prf(average, modifier, msg_start, len(result))
```



LR = 1e-6

model = BertClassifier4(64, 32, 16)

```
[ ] run_model(model, train_dataloader, val_dataloader, LR)
```

```
0%|          | 0/563 [00:00<?, ?it/s]/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2
100%|██████████| 563/563 [06:06<00:00, 1.53it/s]
100%|██████████| 63/63 [00:23<00:00, 2.70it/s]
Epochs: 1 | Train Loss: 0.011 | Train Accuracy: 0.517 | Val Loss: 0.011 | Val Accuracy: 0.630
100%|██████████| 563/563 [06:06<00:00, 1.54it/s]
100%|██████████| 63/63 [00:23<00:00, 2.71it/s]
Epochs: 2 | Train Loss: 0.010 | Train Accuracy: 0.691 | Val Loss: 0.010 | Val Accuracy: 0.785
(24885, 374.5296521782875, 3139, 39.54024076461792)
```

Show confusion matrix

(Second Best accuracy)

```
accuracy measure: 77.92999999999999 %  
precision measure: 77.97828209900118 %  
recall measure: 77.92999999999999 %  
f1 measure: 77.92386839707602 %
```

